BGPmon:
Administrator's Reference Manual

Colorado State University

June 29, 2012

Contents

List of Figures

List of Tables

1 Introduction

BGP Monitor (BGPmon) is a light-weight, scalable, and extensible system for monitoring BGP routing. BGPmon collects routing data by imitating a real BGP router and peering with other BGP routers. All routing updates received from the peers are converted into a convenient XML format and output to interested clients. In addition, BGPmon can output periodic routing tables and control messages. BGPmon can also label routing data to simplify later analysis.

BGPmon clients receive the data from one or more BGPmon instances and perform a variety of data parsing tasks. Any program that can establish a TCP connection and parse XML can become a client. Some clients archive updates to disk, other clients identify and forward only updates that meet some criteria, and so forth.

1.1 Audience

This document is intended for readers interested in installing, configuring, and using BGPmon. Readers interested in the underlying design philosophy of the BGPmon system should refer to the technical paper[1]. Readers interested in understanding the implementation details of BGPmon or modifying the source code should refer to the detailed technical specification found in [2].

1.2 BGPmon Design Overview

The BGPmon design extends the scalable event driven architecture in [3] to meet the requirements of BGP monitoring. BGPmon is a stream based monitoring system; an architectural overview is shown in Figure 1. BGPmon receives BGP messages from peer routers, performs some optional labeling, converts the data to an XML format, and passes the resulting data to clients via TCP connections. BGPmon clients receive an event stream in real-time from peer routers. It is also possible to read historical event streams from archival sources, for example, an archival source could be obtained from RouteViews. BGPmon uses XML to provide easy extensibility, integrate with common tools, and allow local data annotations.


Figure 1: An overview of the BGPmon architecture.

1.2.1 Peers

BGPmon collects BGP updates by establishing peering sessions with routers. BGPmon emulates the behavior of a peering router, but only listens; BGPmon does not implement any routing polices, perform any route

4

selection, manage forwarding tables, or forward packets. The BGPmon design allocates one thread for each peering session. All messages received from a peer are written into the Peer queue for processing by later modules.

1.2.2 Labeling

In the BGPmon design, nearly all data processing is assumed to occur at the clients. However, BGPmon can be configured to add labels to each update received from a BGP peer router. The labels produced by BGPmon include:

NANN	a New ANNouncement (prefix not previously reachable)
DPATH	an update announcing a Different AS PATH
SPATH	an update announcing the Same AS PATH, but some other change
DANN	a Duplicate ANNouncement (no change in any attribute)
WITH	a WITHdraw (prefix no longer reachable)
DWITH	a Duplicate WITHdraw (unreachable prefix withdrawn)

Table 1: The labels assigned to updates by BGPmon.

Calculating the labels requires BGPmon to store RIB_IN tables from each peer. Storing RIB_IN tables can consume substantial amounts of memory and add some computational costs. Therefore, RIB_IN tables can be disabled on a per-peer basis, but labeling is not possible if RIB_IN tables are disabled. Labeling is discussed further in Section 3.7.3.

1.2.3 Chains

A single BGPmon instance can typically handle a large number of peers and clients. However, it may be useful to run multiple instances of BGPmon and connect them together into a single chain. Data from one BGPmon can be fed into a second BGPmon instance. Clients are oblivious to chains and, to the client, it appears as though all data is being collected by a single BGPmon instance. BGPmon chains provide a range of powerful options for scaling to vast numbers of peers, adding additional robustness, or separating peering and collector functions.

As an example, suppose an administrator in Denver wanted to monitor BGP routers at exchange points in Los Angeles and London. A single BGPmon instance running in Denver could peer with routers at both the Los Angeles and London exchange points. However, this requires the use of multi-hop BGP and if the BGPmon instance fails, all data is lost. An alternate strategy would be to deploy a BGPmon instance in London and second BGPmon instance in Los Angeles. The London and Los Angeles BGPmon instances could chain to a third BGPmon instance in Denver. The LA and London BGPmon instances feed data to the Denver instance which in turns provides data to clients. Clients are unaware if the data received comes from a single instance or chain of BGPmons. Section 3.6 describes how to configure BGPmon chains.

1.2.4 Clients

The BGPmon server is designed to simply collect, aggregate and report data; BGPmon itself does not archive data or perform complex analysis on it. These tasks are reserved for BGPmon clients. A client receives XML data from BGPmon via a TCP connection and then performs the desired data processing tasks. Users can build their own clients to suit their individual needs.

1.3 BGPmon Configuration

BGPmon is designed to imitate a router, and its configuration is similar to that of major router vendors. An administrator logs into BGPmon and, after passing authentication checks, has access to the administrative

portion of BGPmon. Similar to the BGP routers sold by large vendors, there are two access levels. Initially, an administrator is connected in *access mode*. Access mode allows the administrator to view statistics, show routing tables, and generally view (but not change) configuration parameters. BGPmon may be configured to allow arbitrary users to login and display statistics, similar to what a user could do if granted access to a router at an ISP.

In order to change the configuration, the user must switch to *enable mode*. An administrator that has entered *enable mode* can perform all BGPmon configuration actions such as adding, deleting, or modifying BGP peers, disabling clients, setting access control policies, and so forth. All changes are stored in memory and will be lost if BGPmon restarts. At any time, an administrator in *enable mode* can save the current BGPmon configuration to a file. The configuration file can then be loaded in subsequent restarts.

If no configuration file is specified at boot time, BGPmon attempts to load a default configuration file. If no configuration is specified and no default configuration file is found, BGPmon allows the administrator to login on port 50000. An alternate default port number can also be specified as a command line parameter to BGPmon.

2   Installing BGPmon

2.1   System Requirements

The BGPmon server does not require specialized resources and is a designed to run on an off-the-shelf Unix box. The processing requirements are minimal and typical boxes can support large numbers of peer connections.

BGPmon clients connect to a BGPmon server via TCP and thus can run on the same machine or on remote machines. Any program that can establish a TCP connection can become a client.

In a recommended configuration, a single box runs the BGPmon server along with a data archiving client to store BGP data for later analysis. Any number of additional clients can connect to BGPmon and receive BGP data in real-time. Adding additional BGP peers and/or clients increases the resource requirements.

2.1.1   Memory Requirements

In its simplest configuration, BGPmon has very little mandatory state and thus can operate with a very small memory footprint A small amount of state is kept for each peering session. As messages arrive from a peer, they are temporarily stored in internal buffers as the messages move from the peering router to the TCP output to clients. Even with hundreds of peers, BGPmon requires only a few megabytes of memory.

BGPmon memory requirements can increase dramatically when routing table (RIB_IN) storage is enabled. As updates are received from a peer, BGPmon can optionally keep track of the peer's current routes in a RIB_IN table. The RIB_IN tables are the primary memory requirements associated with BGPmon. As a general rule of thumb, a peer announcing 250,000 routes will require a 25 MB RIB_IN table

Disabling RIB_IN storage for a peer decreases memory requires, but also prevents BGPmon from labeling updates received from that peer. By comparing an update to the current RIB_IN table, the BGPmon server determines whether the update is a new announcement, path change, duplicate update and so forth as discussed above. This labeling is only possible if RIB_IN storage is enabled for the peer.

The RIB_IN may be used to periodically report the table to clients if the peer does not support BGP route refresh. In most cases, BGPmon periodically re-announces its peers' routing tables to clients by requesting a route refresh from the peer. If the peer does not support route refresh, BGPmon attempts to imitate a route refresh by reporting the RIB_IN file. Periodic routing tables will not be available if the peer router does not support route refresh and BGPmon has been configured not to store RIB_IN tables for the peer.

2.1.2   Operating System Requirements

BGPmon and its clients were developed on Ubuntu and Fedora systems with the objective of being platform independent. Ports to other Unix operating systems are encouraged and some limited help is available from

the BGPmon team. The 7.2.2 release of BGPmon has been testing on Ubuntu 10 and Fedora 14.

2.2 Installation

To perform a default installtion of BGPmon, use:

```
> ./configure
> make
> sudo make install
```

This is the preferred installation procedure for the BGPmon server. A bgpmon account is created and startup scripts are installed under /etc/init.d.

By default, the BGPmon server is installed in *usr/local/bin/bgpmon*. Defaults can be changed by setting the `--prefix` option. Other standard options are available, run `./configure --help` to read more.

=> The libxml2 development library is required and configure will fail in its absence.
Ubuntu:
apt-get install libxml2-dev
Fedora:
yum install libxml2-devel

It is also possible to install BGPmon using a non-root user. To do this one must first disable the creation of the bgpmon user as well as the installation of the startup scripts in init.d.

```
> ./configure --prefix=install/dir --disable-bgpmonuser --disable-initscripts
> make
> make install
```

2.3 Launching BGPmon

To start BGPmon after a root install, simply launch the server using:

```
> sudo service bgpmon start
```

If the installation was not done using the root user the executable should be invoked directly.

```
> install_dir/bin/bgpmon -s -d
```

If this is the first time you are using BGPmon, it is using a default configuration file and the server is waiting for an administrator to login on port 50000. The login port can also be set using command line arguments described in section 2.3.1.

Using the steps discussed in Section 3, an administrator can login and add BGP peers, set access control rules for both future configuration and client access, create BGPmon chains, and apply a variety of optional settings. At a minimum, an administrator will need to configure BGPmon to receive data from at least one BGP peer router, BGPmon chain, or Quagga route collector and allow BGPmon to provide data to at least one client.

7

2.3.1 Optional BGPmon Command Line Arguments

Nearly all configuration is done by logging into BGPmon as discussed in Section 3. However, BGPmon has several optional command line arguments that can useful for some scenarios. Each of these optional values is discussed below.

- $-c < configuration filename >$: Default is $prefix/etc/bgpmon_config.txt. Provides the name of a configuration file to load. The configuration file provides essential information such as the peers to monitor, client access control, and so forth. If no configuration file is specified, BGPmon attempts to load a default configuration file. If the configuration file is not found, BGPmon simply waits for an administrator to login and configure BGPmon.

  Most users will not need to use the *-c filename* option. Unless you specified otherwise, saving a configuration creates the file bgpmon_config.txt and this file is loaded by default when BGPmon restarts. If you plan to have only one BGPmon configuration file, the default configuration file name of bgpmon_config.txt is *strongly recommended*. A site with multiple, distinct configurations may wish to use other file names and can specify which of the multiple configuration files to load using the *-c filename* option.

- $-r < recoveryport >$: Default is 50000. Instructs BGPmon to allow administrator login on the specified port. If this option is not specified, BGPmon uses the port specified in the configuration file or port 50000 if no configuration file is found. Section 3.2.1 describes how to set the login-listener port and save the setting in the configuration file.

  The *-r < recoveryport >* option is intended as a temporary bypass in case either 1) this is the first time BGPmon is running it needs to allow login on a port other than 50000 or 2) the login port set in the configuration file is no longer valid and must be overridden. **The -r recovery-port option takes precedence over any login port found in Configuration File.**

- *-d*: Default is disabled. This option enables daemon mode. This is useful in the case that a non-root install is done and the user wishes to run BGPmon in the background.

- *-s* : Default is that syslog is disabled. This option enables syslog mode and sends all messages to the syslog facility. Syslog mode is recommended and syslog settings can be used to direct BGPmon output to specific file, control the level of output, and so forth.

  A site administrator may modify the source code in order to change the default setting to enable syslog, see [2] for instructions on modifying the default settings.

- $-l < loglevel >$ Default is 4. This option specifies the log level and uses the standard syslog values as follows. Emergencies, Alerts, Critical Errors, and Errors are levels 0 to 3 (respectively). These messages are always logged regardless of the log level setting. Warnings, Notices, Information, and Debug output are levels 4 to 7 (respectively). Setting $loglevel = L$ will log all messages at and below the L. For example, a log level of 4 will display Alerts, Critical Errors, Errors, and Warnings, but will not display Notices, Information, or Debugging output. A site administrator may modify the source code in order to change the default logvalue, see [2] for instructions on modifying the default settings.

- $-f < facility >$ Default is USER. option specifies the syslog Facility. A site administrator may modify the source code in order to change the default logvalue, see [2] for instructions on modifying the default settings.

3   Configuring BGPmon

The Command Line Interface (CLI) is used to configure BGPmon. Once connected, a user is initially set to *access mode* which allows them to view statistics, show routing tables, and generally view (but not change)

configuration parameters. In order to change the configuration settings, a user must switch to *enable mode*. In this mode a user can perform BGPmon configuration actions such as adding, deleting, or modifying BGP peers and chains, disabling clients, and setting access control policies.

At any time, a user in *enable mode* can save the current BGPmon configuration so it can be loaded the next time BGPmon starts. **Any configuration changes are stored in memory and will be lost if BGPmon restarts. To make changes permanent, the administrator must save the BGPmon configuration.**

This section describes how to configure BGPmon, beginning with logging into BGPmon and proceeding through the steps to configure future login access, enable clients, configure chains, create peers, and finally set optional parameters. A first time administrator should *read each subsection in order* and follow the configuration steps in that section. An experienced administrator may want to skip directly to the relevant subsection. A complete command reference is also available beginning in Appendix A.

3.1 Logging Into BGPmon

If BGPmon has not been previously configured, then the server will be listening on the loopback address and port 50000 for incoming connections. Use telnet to connect to the BGPmon server. After connecting to BGPmon, you will be prompted for a guest password. By default, the guest password is 'BGPmon'.

```
> telnet 127.0.0.1 50000
Trying 127.0.0.1...
Connected to 127.0.0.1.
Escape character is '^]'.
Password:
```

=> loopback addresses
IPv4 loopback address is 127.0.0.1
IPv6 loopback address is ::1

Initially, a user is connected in *access mode* and must change to *enable mode* to alter the server's configuration settings. To enter *enable mode*, use *enable*

```
host> enable
enable password:
```

You will then be prompted for the enable password, which is also 'BGPmon' by default. From *privileged mode* you need to enter *configure mode* to gain access to all the configuration commands. To enter this mode type:

```
host# configure
host(config)#
```

You are now ready to configure BGPmon. It is recommended to change the default passwords and configure restrictions on future login accesses as discussed in Section 3.3.

In the future, if you are the BGPmon administrator and have lost the login settings, Section 3.1.1 describes how to reset these values without losing the other configuration information.

3.1.1 BGPmon Login Recovery

Due to configuration errors and/or misplaced information, an administrator may be unable to login to BGPmon. For example, the administrator may have lost the access password or set up an overly aggressive access control list to the BGPmon Command Line Interface. In these cases, the administrator has four options for recovery.

The first option is to use the *-r recovery-port* parameter described in section 2.3. This parameter will override the login-listener's port so an administrator can still use the Command Line Interface if there is a conflict between BGPmon and another application trying to use the same port. However, the recovery-port will not allow an administrator to bypass the ACL specified for the login-listener. If this is the problem, then one of the following options for recovery must be used.

The second option can be used in cases where the configuration is simple or very close to the default settings. Simply delete the active configuration, which is stored in 'bgpmon_config.txt', and let BGPmon start with the default settings. Then make any necessary changes to the BGPmon configuration and save the settings.

There are many times where starting over or simply bypassing the default port won't be a viable option and for these cases there are other options. The first is trying to recover an old configuration. Any configuration changes that are saved will be written into 'bgpmon_config.txt'. Every time a new configuration is saved the old configuration is archived in the same directory with the time and date preceding the file name. An example archived configuration file is '1041_2252009_bgpmon_config.txt', which was created at 10:41am on 02/25/2009. To find a suitable recovery point, start by backing up the current configuration then copying the archived configuration over the current configuration. Hopefully within a few tries a good recovery point can be identified.

The final option is to manually edit the active configuration with a text editor. It's easy to manually edit the active configuration for BGPmon since all settings are stored as XML in 'bgpmon_config.txt'. By comparing the last known good configuration and the current configuration all changes between these versions can be identified then reapplied to the current configuration. Just be careful to make changes slowly until the unwanted change to BGPmon is identified and corrected.

3.2 Saving Configuration Changes

Any configuration changes made to BGPmon through the CLI will be lost unless the user saves the configuration file.

```
host(config)# exit
host# copy running-config startup-config
```

The above command will copy the in-memory configuration (or the running-config) to the saved configuration file (startup-config).

3.3 Configuring Login Access To BGPmon

Login access is controlled in BGPmon by the login listener. This module will bind to an address and port then listen for incoming Command Line Interface connections. When a connection is established it will check to see if the address of the new connection is allowed or disallowed based on the Access Control List for the Login Listener. Use the following commands to see what the current settings of the login-listener are:

```
host> show login-listener address
login-listener address is 127.0.0.1
host> show login-listener port
login-listener socket is 50000
```

```
host> show login-listener acl
Current acl: permitall
```

3.3.1   Configuring the login-listener

=> In order to perform configuration the user must be in config mode

There are three main components that can be configured for login-listener. First is the address that the login-listener will attempt to use when BGPmon starts. To change the login-listener's address, type the following command:

```
host(config)# login-listener address <new-address>
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check the *new-address* to make sure that it's a valid IP address and is available on the local machine. If the address isn't valid then a warning will be returned and nothing is set. Also, any loopback address in Appendix B can be used.

The next major component to configure is the port. To change the port, type the following command:

```
host(config)# login-listener port <new-port>
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check to make sure the *new-port* is a valid port number but will not check to see if the port can be bound. So, if the port for the login-listener is set to an unavailable port and BGPmon is restarted then the Command Line Interface will not be available upon restart. If this happens see section 3.1.1 about recovering.

The final component for the login-listener is the Access Control List, which is the list that controls which addresses are allowed to connect or not connect. To set the active ACL, type the following:

```
host(config)# login-listener acl <acl-name>
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check to make sure the *acl-name* is valid within BGPmon before setting it. Refer to section 3.5 to learn about configuring Access Control Lists.

3.4   Configuring Client Access

Similar to the login-listener, client access is controlled in BGPmon by the client-listener. This module binds to two addresses and ports and listens for incoming connections. One of the ports listens for clients requesting BGP update data and the other listens for requests for RIB table data. When a connection is established it will check to see if the address of the new connection is allowed or disallowed based on rules setup in the Access Control List. Use the following commands to see what the current settings of the client-listener are:

```
host> show client-listener status
client-listener is enabled
```

```
host> show client-listener summary
Client-listener is enabled

UPDATE ACL: permitall
UPDATE address is 127.0.0.1
UPDATE port is 50001

RIB ACL: permitall
RIB address is 127.0.0.1
RIB port is 50002
host> show client-listener update ?
 - acl
 - port
 - address
host> show client-listener update acl
acl: permitall
host> show client-listener update port
port is 50001
host> show client-listener update address
address is 127.0.0.1
host> show client-listener rib ?
 - acl
 - port
 - address
host> show client-listener rib acl
acl: permitall
host> show client-listener rib port
port is 50002
host> show client-listener rib address
address is 127.0.0.1
```

3.4.1 Configuring the client-listener

=> In order to perform configuration the user must be in config mode

There are four main components that can be configured for the client-listener. Any change made to these components will result in the client-listener stopping then starting, if necessary, with the new values. The first component are the addresses that the client-listener will use. To change the addresses type the following commands:

```
host(config)# client-listener update address <new_update_address>
host(config)# client-listener rib address <new_rib_address>
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check the *new-address* to make sure that it's a valid IP address and is available on the local machine. Also, any loopback address in Appendix B can be used.

The next major component to configure is the port. To change the port, type the following command:

```
host(config)#client-listener update port <new_update_port>
host(config)#client-listener rib port <new_rib_port>
```

```
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check to make sure the *new-port* is a valid port number but will not check to see if the port can be bound until the client-listener stops then attempts to start again.

The next component for the client-listener is the Access Control List, which is the list that controls which addresses are allowed to connect or not connect. To set the active ACL, type the following:

```
host(config)# client-listener update acl <acl_name>
host(config)# client-listener rib acl <acl_name>
host(config)# exit
host# copy running-config startup-config
```

BGPmon will check to make sure the *acl-name* is valid within BGPmon before setting it. Refer to section 3.5 to learn about configuring Access Control Lists.

The final component is the status, which can be set to either *enabled* or *disabled*. Use the following commands to change the status:

```
host(config)# client-listener <disable | enable>
```

3.5   Access Control Lists

Access Controls Lists are used within BGPmon to control which addresses are or aren't allowed to connect. BGPmon initially comes with two ACLs: 'denyall' which denies all traffic and 'permitall' which permits all traffic. To see a list of ACLs that are available, type *show acl*. The optional parameter *acl_name* may be used to limit the results to a specific ACL.

```
host> show acl
ACL name:denyall
index address               mask                type
0     any                   any                 deny

ACL name:permitall
index address               mask                type
0     any                   any                 permit

host> show acl denyall
ACL name:denyall
index address               mask                type
0     any                   any                 deny

```

3.5.1   ACL inverse mask logic

Each ACL is made up of a series of rules and each rule is made up of three components. The first two components, address and mask, are used to determine whether the rule is applicable to the incoming address. The third component, type, tells BGPmon whether to permit or deny access to addresses that satisfy the rule.

Tables 2 and 3 are examples of how BGPmon determines if a rule is applicable to incoming address. The first step is to OR the rule's address and mask together then OR the incoming address and mask together. Subtracting these two results will indicate whether a rules is applied. A zero value means that the rule is applied and a non-zero value means that the rule does not apply.

Table 2: An ACL example resulting in the application of the rule.

Example 1		
	String	**Binary**
Incoming address	10.1.1.255	00001010 00000001 00000001 11111111
Rule address	10.1.1.1	00001010 00000001 00000001 00000001
Rule mask	0.0.255.255	00000000 00000000 11111111 11111111
incoming address | mask		00001010 00000001 11111111 11111111
rule address | mask		00001010 00000001 11111111 11111111
zero difference - rule should be applied		00000000 00000000 00000000 00000000

Table 3: An ACL example resulting in the rule not being applied.

Example 2		
	String	**Binary**
Incoming address	10.1.255.255	00001010 00000001 11111111 11111111
Rule address	10.1.1.1	00001010 00000001 00000001 00000001
Rule mask	0.0.0.255	00000000 00000000 00000000 11111111
incoming address | mask		00001010 00000001 11111111 11111111
rule address | mask		00001010 00000001 00000001 11111111
non-zero difference - rule should not be applied		00000000 00000000 11111110 00000000

3.5.2 Creating and Editing ACLs

=> In order to perform configuration the user must be in config mode

To create an ACL each of the three components, along with a name must be entered. The first two components, address and mask, are used to determine whether the rule is applicable to the incoming address. The third component tells BGPmon whether to permit or deny applicable addresses. The following is an example that creates an ACL with the same functionality as "permit all."

```
host(config)#acl acl-test
Created ACL, now editing: acl-test
host(config-acl)#permit any 0
host(config-acl)#
```

The first step to creating or editing an ACL is to begin an editing session associated with that ACL.

```
host(config)#acl <acl_name>
```

This command will attempt to open the specified ACL. If it is not found, then a new ACL is created. Within the ACL edit mode there are several commands that allow the user to maintain an ACL.

The first set of commands are used to create 'permit' or 'deny' rules.

```
host(config-acl)#permit any [rule_index]
host(config-acl)#permit address mask [rule_index]
host(config-acl)#deny any [rule_index]
host(config-acl)#deny address mask [rule_index]
```

In each of these commands either the 'permit' or 'deny' keyword is specified then followed by a series of rules that indicates the range of addresses the rule applies to. The parameter *address* is the IP address used in the rule. The parameter *mask* is used to specify which bits in the address are significant and which should be ignored. If a rule should be applied to all addresses then use the keyword 'any' instead of the address/mask pair. Finally the optional parameter *rule_index*, if set, will be used to specify where in the list the rule should be inserted. When a rule is inserted at a rule_index, then all rules that follow will be incremented by one.

To remove a rule from the list use the following command:

```
host(config-acl)#no <rule_index>
```

When a rule is removed from the list all other rules in the list will be re-indexed while maintaining their relative ordering.

3.5.3   Deleting ACLs

To delete an ACL and all associated rules use the following command:

```
host(config)#no acl <acl_name>
```

It is important to note that when an ACL is removed, any modules using that ACL will be set to the default behavior, which is to deny all traffic.

3.6   Configuring Chains

Chaining is used in BGPmon to connect multiple instances of BGPmon together so XML messages from one BGPmon instance can be sent to another BGPmon instance. This allows for easy distribution of data collection from peers at many distant locations.

3.6.1   Creating and updating a chain

To create a chain use the following command:

```
host(config)#chain address <[update:port] [rib:port]> [retry:retry_interval]
```

The only required parameters are address and port. Any parameter which is not specified will have a default value assigned to it. By default, the *retry_interval* is set to 60, and the chain is enabled. If the user wishes to manually shut off or restore a chain, use the command:

```
host(config)#chain address <[update:port] [rib:port]> <[enable] [disable]>
```

Once a chain has been created, the same command can be used to update settings for that chain. For example, assume a chain was created with the following command:

```
host(config)# chain 192.168.1.1 update:50001
```

To set the status of this chain to disabled and retry interval to 10, use the following command:

```
host(config)# chain 192.168.1.1 disable retry:10
```

One important thing to remember about creating chains is that the address and port uniquely identify a chain. So, assume a chain is created with the following command:

```
host(config)# chain 192.168.1.1 update:50001
```

Now, assume this command is run:

```
host(config)# chain 192.168.1.1 rib:50002
```

The second command will create a new chain at *192.168.1.1* on port *50002* while the first chain will remain at *192.168.1.1* on port *50001*. **Once a chain has been created its address and port cannot updated.** Any command attempting to do this will create another chain.

3.6.2 Delete a chain

A chain can be deleted with the following command:

```
host(config)#no chain address [port]
```

A port doesn't need to be specified if the default port was used when creating the chain. When this command is issued BGPmon will set the chain to disabled then then mark the chain to be deleted.

3.7 Configuring Peers and Peer Groups

A Peer is router that BGPmon maintains a connection with and every Peer has its own set of configurable parameters. Peer Groups are a way of sharing configurations between sets of Peers .

The first step in configuring a peer is to enter the router configuration mode. To do this, enter the following command from the configuration mode prompt:

```
host(config)#router bgp <AS_Number>
```

The *AS_Number* entered will be used as the local AS number for any peering sessions that are created. Every command associated with a peer or peer-group in this mode follows the same structure: a base command used to identify the peer or peer-group followed by a command to modify that peer or peer-group. The base of the command is as follows:

```
host(config-router)#neighbor <address | peer-group> [port <port>]
```

If a valid IP address is specified for the *address* then the command will refer to a peer. The optional *port* parameter can then be used to specify a custom port, otherwise the default port 179 will be used. If the *address* is not a valid IP address then the command will refer to a peer-group.

3.7.1 Creating Peers

To create a peer, run any command where the base command has a valid address. If the peer does not exist then BGPmon will attempt to create the peer. Example:

```
host(config-router)#neighbor 192.168.1.1 remote as 9552
```

If no pre-existing peers are configured, then this command would create a peer with an IPv4 address of *192.168.1.1* and an AS number of 9552.

To create a peer with IPv6 address, run following commands:

=> Note: this is a workaround for a known issue in version 7.2.3. The procedure for peering with an IPv6 router will change dramatically in the future.

```
host(config-router)# neighbor fd68:e916:5287:9f27::7 remote as 3552
host(config-router)# neighbor fd68:e916:5287:9f27::7 local bgpid 129.82.138.29
host(config-router)# exit
host(config)# exit
host# copy running-config startup-config
host# exit
host> exit
```

Shutdown BGPmon server:

```
> sudo service bgpmon stop
```

Open */usr/local/etc/bgpmon_config.txt* with your favorite text editor. Under the *<PEER>* tag that contains your IPv6 router information, add the following line:

```
<MONITOR_IP_ADDR>ipv6any</MONITOR_IP_ADDR>
```

Save configuration and restart BGPmon server:

```
> sudo service bgpmon start
```

If no pre-existing peers are configured, then this command would create a peer with an IPv6 address of *fd68:e916:5287:9f27::7* and an AS number of 3552. In this example, BGPmon uses the BGPID 129.82.138.29.

3.7.2 Creating Peer-groups

To create a peer-group the following command is used, where *peer_group_name* will be the name of the peer-group and not a valid IP address:

```
host(config-router)#neighbor <peer_group_name> peer-group
```

Once a peer-group has been created then peers can be assigned to the group with the following command:

```
host(config-router)#neighbor <address> [port <port>] peer-group <peer_group_name>
```

A peer can be moved to another peer-group with the same command and use the following command to remove a peer from a peer-group:

```
host(config-router)#no neighbor <address> [port <port>] peer-group <peer_group_name>
```

3.7.3 Configuring Parameters

Within a Peer or Peer-group there are many different parameters that can be configured. The first group are the commands that configure parameters for the *remote* and *local* BGP versions. These commands are as follows (the command prompt is omitted for space concerns):

```
neighbor <[address] [peer-group]> [port <port>] <[remote] [local]> as <as_number>
neighbor <[address] [peer-group]> [port <port>] <[remote] [local]> bgpid <bgpid>
neighbor <[address] [peer-group]> [port <port>] <[remote] [local]> bgp-version <bgp_version>
neighbor <[address] [peer-group]> [port <port>] <[remote] [local]> hold-time <hold_time>
```

The second group of commands are the labeling commands. These commands control how messages are labeled in each peer:

```
neighbor <[address] [peer-group]> [port <port>] label-action NoAction
neighbor <[address] [peer-group]> [port <port>] label-action Label
neighbor <[address] [peer-group]> [port <port>] label-action StoreRibOnly
```

The final group of commands are commands used to control the actual peer. These commands will trigger a route-refresh, enable and disable a peer, and set the md5 encryption for a password.

```
neighbor <[address] [peer-group]> [port <port>] route-refresh
neighbor <[address] [peer-group]> [port <port>] enable
neighbor <[address] [peer-group]> [port <port>] disable
neighbor <[address] [peer-group]> [port <port>] md5 <md5_password>
```

3.7.4 Configuring Capabilities

For every peer there are a set of *announce* and *receive* capabilities that can be configured. Announce capabilities are those that will be included in the BGP open message when establishing a peering connection. Receive capabilities are configured to require, allow or forbid capabilities announced in the open message from the peer. By default any capability not mentioned as a receive capability defaults to allowed.

The three most common capabilities to include are 4-byte ASNs, IPv4 Multiprotocol and IPv6 Multiprotocol. These capabilities must be configured manually in the configuration file. Before starting BGPmon you can configure these capabilities by adding the following XML to the configuration file.

```
<PEERS>
  <PEER>
    <MONITOR_PORT>[bgpmon's monitoring port]</MONITOR_PORT>
    <MONITOR_AS>[bgpmon's AS]</MONITOR_AS>
    <PEER_IP_ADDR>[the peer's IP address]</PEER_IP_ADDR>
    <PEER_PORT>[The peer's port (usually 179]</PEER_PORT>
    <PEER_AS>[The peer's ASN]</PEER_AS>
    <ANNOUNCE_CAPABILITY_LIST>
      <!-- 4-byte ASN capability -->
      <CAPABILITY>
        <CODE>65</CODE>
        <LENGTH>4</LENGTH>
        <VALUE>[BGPmon's 4 byte ASN in HEX]</VALUE>
      </CAPABILITY>
      <!-- IPv4 multi-protocol capability -->
      <CAPABILITY>
        <CODE>1</CODE>
        <LENGTH>4</LENGTH>
        <VALUE>00010001</VALUE>
      </CAPABILITY>
      <!-- IPv6 multi-protocol capability -->
      <CAPABILITY>
        <CODE>1</CODE>
        <LENGTH>4</LENGTH>
        <VALUE>00020001</VALUE>
      </CAPABILITY>
    </ANNOUNCE_CAPABILITY_LIST>
  </PEER>
</PEERS>
```

=> Note: When expressing the ASN in hex there is no need for a prefix of 0x. For example, an ASN of 100004 would be expressed as 000186a4. In the case that you are announcing the capability of handling 4 bytes ASNs, but have a 2 byte ASN simply express the 2 byte ASN in the same space prepended with zeros to take up 4 bytes.

3.7.5 Deleting a peer

To delete a peer, use the following command:

```
host(config-router)#no neighbor <address> [port <port>]
```

To delete a peer-group, use the following command:

```
host(config-router)#no neighbor <peer-group>
```

3.7.6 Resetting a peer's connection

Resetting a peer will close the peer's connection then attempt to reopen it. Any new settings that have been applied to the peer will be applied when the peer starts again. To reset a connection to a peer use the following command:

```
host(config-router)#clear neighbor <address> [port <port>]
```

A   Command Line Interface Reference

Privilege Level: Access Control List Configuration
Mode commands:
exit - Sets the security level of the current user back to Configure.
end - Sets the security level of the current user back to Privileged.
Rule commands:
show acl [*acl_name*] - Show the rules for all [or a single] ACL.
no <*rule_number*> - Deletes an ACL rule based on the rule number.
no permit <*rule_number*> - Deletes a permit ACL rule based on the rule number.
no deny <*rule_number*> - Deletes a deny ACL rule based on the rule number.
no label <*rule_number*> - Deletes a deny ACL rule based on the rule number.
no ribonly <*rule_number*> - Deletes a deny ACL rule based on the rule number.
permit any [*rule_index*] - Creates a rule that permits any address. If this rule is applied to the MRT module then only updates are allowed from the address.
deny any [*rule_index*] - Creates a rule that denies any address.
label any [*rule_index*] - Creates a rule that allows only LABEL messages. This rule should only be applied to the MRT module.
ribonly any [*rule_index*] - Creates a rule that allows only RIB messages. This rule should only be applied to the MRT module.
permit <*address*><*mask*>[*rule_index*] - Creates a rule that permits any address which matches the address and mask combination specified in the rule. If this rule is applied to the MRT module then only updates are allowed from the address.
deny <*address*><*mask*>[*rule_index*] - Creates a rule that denies any address which matches the address and mask combination specified in the rule.
label <*address*><*mask*>[*rule_index*] - Creates a rule that allows LABEL messages from an address which matches the address and mask combination specified in the rule.
ribonly <*address*><*mask*>[*rule_index*] - Creates a rule that allows only RIB messages any address which matches the address and mask combination specified in the rule.

Privilege Level: Router Configuration

Announce and Receive commands:

[**no**] **neighbor** <[*peer-group*] [*address*]>[**port** *neighbor_port*] **announce** <*code*><*length*><*value*> - [Removes] Adds a custom announce capability to the neighbor or peer-group specified.

[**no**] **neighbor** <[*peer-group*] [*address*]>[**port** *neighbor_port*] **announce** <[**ipv4**] [**ipv6**]><[**unicast**] [**multicast**]> - [Removes] Adds a predefined announce capability to the neighbor or peer-group specified.

[**no**] **neighbor** <[*peer-group*] [*address*]>[**port** *neighbor_port*] **receive** <[**require**] [**refuse**] [**allow**]><*code*><*length*><*value*> - [Removes] Adds a custom receive capability to the neighbor or peer-group specified.

[**no**] **neighbor** <[*peer-group*] [*address*]>[**port** *neighbor_port*] **receive** <[**require**] [**refuse**] [**allow**]><[**ipv4**] [**ipv6**]><[**multicast**] [**unicast**]> - [Removes] Adds a predefined receive capability to the neighbor or peer-group specified.

Creating a peer-group:

neighbor <*peer_group_name*>**peer-group** - Creates a peer-group.

neighbor <*address*>[**port** *neighbor_port*] **peer-group** *peer_group_name* - Assigns a peer to a peer-group.

Mode commands:

exit - Sets the security level of the current user back to Configure.

end - Sets the security level of the current user back to Privileged.

Remote and Local commands:

neighbor <[**address**] [**peer-group**]>[**port** <**neighbor port**>] <[**remote**] [**local**]>**as** <**as_number**> - Sets either the remote or local AS number for a peer or peer-group.

neighbor <[**address**] [**peer-group**]>[**port** <**neighbor port**>] <[**remote**] [**local**]>**bgpid** <*bgpid*> - Sets either the remote or local BGPID for a peer or peer-group.

neighbor <[**address**] [**peer-group**]>[**port** <**neighbor port**>] <[**remote**] [**local**]>**bgp-version** <*bgp_version*> - Sets either the remote or local BGP version for a peer or peer-group.

neighbor <[**address**] [**peer-group**]>[**port** <**neighbor port**>] <[**remote**] [**local**]>**hold-time** <*hold_time*> - Sets either the remote or local hold-time for a peer or peer-group.

Updating a peer or peer-group:

no neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] - Deletes a neighbor or peer-group.

clear neighbor <*address*>[**port** *neighbor_port*] - Resets the connection with the neighbor.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **enable** - Enables a peer's connection.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **label-action NoAction** - Sets the labelling action for a peer-group or peer to NoAction, stopping all labelling.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **label-action Label** - Sets the labelling action for a peer-group or peer to Label, starting all labelling.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **label-action StoreRibOnly** - Sets the labelling action for a peer-group or peer to StoreRibOnly, stopping labelling but storing the RIB-IN table.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **route-refresh** - Triggers a route-refresh for that peer from the BGPmon RIBIN table.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **md5** [*md5_password*] - Sets the MD5 password used in authenticating with a peer.

neighbor <[**address**] [**peer-group**]>[**port** *neighbor_port*] **disable** - Disables a peer's connection.

Privilege Level: Configuration
ACL commands:
no acl *<acl name>* - Deletes an Access Control List.
acl *<acl_name>* - Creates a new ACL or opens an existing ACL then changes the security level of the user to ACL Configuration.
Chain commands:
chain *<address><*[**update:***port*] [**rib:***port*]*>*[[**enable**] [**disable**]] [**retry:***retry_interval*] - Creates a chain to the given address. If *port* or *retry* are not specified, BGPmon will use its default values. The additional options are not necessary to enable or disable the chain.
no chain *<address>*[*port*] - Deletes a Chain.
Client commands:
client-listener update address *<local-address>* - Sets the address that the client update listener will listen on.
client-listener update port *<port>* - Sets the port that the client update module will listen on for new client connections.
client-listener update acl *<acl_name>* - Sets the ACL which will be applied to the clients update modules.
client-listener rib address *<local-address>* - Sets the address that the client rib-stream listener will listen on.
client-listener rib port *<port>* - Sets the port that the client rib module will listen on for new client connections.
client-listener rib acl *<acl_name>* - Sets the ACL which will be applied to the clients rib modules.
client-listener *<*[**enable**] [**disable**]*>* - Enables or disables the listener, both update and rib, for the client module.
Login commands:
login-listener address *<local-address>* - Sets the address that the login-listener will listen on.
login-listener port *<port>* - Sets the port that the login-listener will listen on.
login-listener acl *<acl_name>* - Sets the ACL that will be applied to the login-listener.
Mode commands:
exit - Sets the security level of the current user to Privileged.
end - Sets the security level of the current user to Privileged.
Neighbor commands:
router bgp *<monitor_AS_number>* - Allows the user to enter Router Configuration mode where peers and peer groups can be modified.
Periodic commands:
periodic route-refresh *<*[**disable**] [**enable**]*>* - Turns off/on periodic route refresh.
periodic route-refresh *<interval>* - Sets the interval for periodic route refresh.
periodic status-message *<status_message_interval>* - Sets the interval for sending periodic status messages.

Privilege Level: Configure
MRT commands:
mrt-listener port *<port_number>* - Sets the port number for the MRT listener.
mrt-listener address *<address>* - Sets the address for the MRT listener. This may be any valid IP address, or one of the keywords in Appendix A.
mrt-listener acl *<acl_name>* - Sets the ACL for the MRT listener.
mrt-listener enable - Enables the MRT listener.
mrt-listener disable - Shuts off the MRT listener.
Queue commands:
queue pacingOnThresh *<pacingOnThresh>* - Modifies the pacing-on threshold value for queues.
queue pacingOffThresh *<pacingOffThresh>* - Modifies the pacing-off threshold value for queues.
queue alpha *<alpha>* - Modifies the alpha value for queues.
queue minWritesLimit *<minimum_write_limit>* - Modifies the minimum write limit for queues.
queue pacingInterval *<pacing_interval>* - Modifies the maximum write limit for queues.
Password Commands:
password *<new_password>* - Sets the login password.
enable password *<new_password>* - Sets the enable-mode password.
Privilege Level: Privileged
Client commands:
kill client *<client_id>* - Disconnects the client specified by the client_id.
Configuration commands:
copy running-config *<filename>* - Copies the running configuration to the specified filename.
copy running-config startup-config - Copies the running configuration to the startup configuration.
shutdown - Shuts down the BGPmon server.
Mode commands:
configure - Sets the security level of the current user to Configure.
disable - Sets the security level of the current user back to Guest.
exit - Sets the security level of the current user back to Guest.

Privilege Level: Guest
ACL commands:
show acl [*acl_name*] - Shows the details of an Access Control List.
Chain commands:
show chains - Shows information about all currently configured chains.
MRT commands:
show mrt neighbor - Shows the connection information for the MRT collectors that peer with BGPmon.
show mrt clients - Shows connection information about the routers that are providing BGP data via an MRT collector.
Client commands:
show clients - Shows the list of users actively connected on the Client module.
show client-listener status - Shows if the client listener is enabled or disabled.
show client-listener update port - Shows the current port for the client-listener's update channel.
show client-listener update address - Shows the current address for the client-listener's update channel.
show client-listener update acl - Shows the Access Control List assigned to the client-listener's update channel.
show client-listener rib port - Shows the current port for the client-listener's rib channel.
show client-listener rib address - Shows the current address for the client-listener's rib channel.
show client-listener rib acl - Shows the Access Control List assigned to the client-listener's rib channel.
show client-listener summary - Shows all of the information about the update and rib client listeners.
Configuration commands:
show running - Shows running configuration for the instance of BGPmon.
Login commands:
show login-listener port - Shows the current port for the login-listener.
show login-listener address - Shows the current address for the login-listener.
show login-listener acl - Shows the Access Control List assigned to the Command Line Interface module.
Mode commands:
enable - Sets the security level of the current user to Privileged.
exit - Exits the BGPmon Command Line Interface.
BGP commands:
show bgp neighbor [*neighbor_address*] [**port** *neighbor_port*] - Shows information about all [or one] configured peers.
show bgp routes [*neighbor_address*] - Shows all routes learned from all [or one] BGP or MRT peers.
show bgp prefix [*prefix*] - Shows information about all [or a single] prefix.
Periodic commands:
show periodic route-refresh - Shows the route-refresh interval.
show periodic route-refresh status - Shows whether route refresh is currently enabled or disabled.
show periodic status-message - Shows the status message interval.

Privilege Level: Guest
MRT commands:

show mrt-listener status - Shows if the mrt listener is enabled or disabled.
 	
     show mrt-listener port - Shows the current port for the mrt-listener.
 	
     show mrt-listener address - Shows the current address for the mrt-listener.
 	
     show mrt-listener acl - Shows the Access Control List assigned to the mrt-listener.
 	
     show mrt-listener summary - Shows all of the other commands' output at once.

Queue commands:
show queue [*queueName*] - Shows information about all [or one of] BGPmon's internal queues.

B   IP Address reference

Shortcut	Version	Actual
ipv4any	IPv4	any address
ipv6any	IPv6	any address
ipv4loopback	IPv4	loopback address = 127.0.0.1
ipv6loopback	IPv6	loopback address = ::1

References

[1] D. Matthews, N. Parrish, H. Yan, , and D. Massey. BGPmon: A real-time, scalable, extensible monitoring system. *Proceedings of the ACM SIGCOMM Internet Measurement Confernce (IMC)*, 2008.

[2] D. Matthews, H. Yan, , and D. Massey. Bgpmon implementation and technical specification, 2008.

[3] M. Welsh, D. Culler, and E. Brewer. Seda: an architecture for well-conditioned, scalable internet services. In *SOSP '01: Proceedings of the eighteenth ACM symposium on Operating systems principles*, pages 230–243, New York, NY, USA, 2001. ACM.